

LOGISTIC REGRESSION

```
In [2]: #Name : Shreya Sharma
#Roll no. : 46
#Sectin : 3B
#Date : 05/10/2024
```

```
In [3]: #Aim : To perform operation on logestic regression
```

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

```
In [4]: import os
```

```
In [5]: os.getcwd()
```

```
Out[5]: 'C:\\Users\\pravi'
```

```
In [6]: os.chdir("C:\\Users\\pravi\\Desktop")
```

```
In [11]: df=pd.read_csv("framingham.csv")
```

```
In [13]: df.head()
```

```
Out[13]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	
0	1	39	4.0	0	0.0	0.0	0	0	0	195.0	106.0	70.0	26
1	0	46	2.0	0	0.0	0.0	0	0	0	250.0	121.0	81.0	28
2	1	48	1.0	1	20.0	0.0	0	0	0	245.0	127.5	80.0	28
3	0	61	3.0	1	30.0	0.0	0	1	0	225.0	150.0	95.0	28
4	0	46	3.0	1	23.0	0.0	0	0	0	285.0	130.0	84.0	28

```
In [15]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4238 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4238 non-null  int64
1   age                   4238 non-null  int64
2   education             4133 non-null  float64
3   currentSmoker         4238 non-null  int64
4   cigsPerDay            4209 non-null  float64
5   BPMeds               4185 non-null  float64
6   prevalentStroke       4238 non-null  int64
7   prevalentHyp         4238 non-null  int64
8   diabetes              4238 non-null  int64
9   totChol              4188 non-null  float64
10  sysBP                4238 non-null  float64
11  diaBP                4238 non-null  float64
12  BMI                  4219 non-null  float64
13  heartRate            4237 non-null  float64
14  glucose              3850 non-null  float64
15  TenYearCHD           4238 non-null  int64
dtypes: float64(9), int64(7)
memory usage: 529.9 KB
```

```
In [17]: df.isna().sum()
```


```
Out[17]: male          0
age          0
education    105
currentSmoker 0
cigsPerDay   29
BPMeds       53
prevalentStroke 0
prevalentHyp 0
diabetes     0
totChol      50
sysBP        0
diaBP        0
BMI          19
heartRate    1
glucose      388
TenYearCHD   0
dtype: int64
```

```
In [19]: df
```

Out[19]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP
0	1	39	4.0	0	0.0	0.0	0	0	0	195.0	106.0	70.0
1	0	46	2.0	0	0.0	0.0	0	0	0	250.0	121.0	81.0
2	1	48	1.0	1	20.0	0.0	0	0	0	245.0	127.5	80.0
3	0	61	3.0	1	30.0	0.0	0	1	0	225.0	150.0	95.0
4	0	46	3.0	1	23.0	0.0	0	0	0	285.0	130.0	84.0
...
4233	1	50	1.0	1	1.0	0.0	0	1	0	313.0	179.0	92.0
4234	1	51	3.0	1	43.0	0.0	0	0	0	207.0	126.5	80.0
4235	0	48	2.0	1	20.0	NaN	0	0	0	248.0	131.0	72.0
4236	0	44	1.0	1	15.0	0.0	0	0	0	210.0	126.5	87.0
4237	0	52	2.0	0	0.0	0.0	0	0	0	269.0	133.5	83.0

4238 rows × 16 columns



```
In [21]: df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)
```

```
In [23]: df['education'].fillna(value = df['education'].mean(),inplace=True)
```

```
In [25]: df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)
```

```
In [27]: df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)
```

```
In [29]: df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)
```

```
In [31]: df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)
```

```
In [33]: df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)
```

```
In [35]: df.isna().sum()
```

```
Out[35]: male          0
age          0
education     0
currentSmoker 0
cigsPerDay    0
BPMeds        0
prevalentStroke 0
prevalentHyp  0
diabetes      0
totChol       0
sysBP         0
diaBP         0
BMI           0
heartRate     0
glucose       0
TenYearCHD    0
dtype: int64
```

```
In [37]: x = df.drop("TenYearCHD",axis=1)
y = df['TenYearCHD']
```

```
In [39]: x
Out[39]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP
0	1	39	4.0	0	0.0	0.00000	0	0	0	195.0	106.0	70.0
1	0	46	2.0	0	0.0	0.00000	0	0	0	250.0	121.0	81.0
2	1	48	1.0	1	20.0	0.00000	0	0	0	245.0	127.5	80.0
3	0	61	3.0	1	30.0	0.00000	0	1	0	225.0	150.0	95.0
4	0	46	3.0	1	23.0	0.00000	0	0	0	285.0	130.0	84.0
...
4233	1	50	1.0	1	1.0	0.00000	0	1	0	313.0	179.0	92.0
4234	1	51	3.0	1	43.0	0.00000	0	0	0	207.0	126.5	80.0
4235	0	48	2.0	1	20.0	0.02963	0	0	0	248.0	131.0	72.0
4236	0	44	1.0	1	15.0	0.00000	0	0	0	210.0	126.5	87.0
4237	0	52	2.0	0	0.0	0.00000	0	0	0	269.0	133.5	83.0

4238 rows × 15 columns

```
In [41]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=42)
```

```
In [43]: x_train
```

```
Out[43]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP
3252	1	40	4.0	1	30.0	0.0	0	0	0	205.0	131.0	81.0
3946	0	57	2.0	0	0.0	0.0	0	1	0	250.0	152.5	92.5
1261	0	47	1.0	0	0.0	0.0	0	0	0	230.0	123.0	71.0
2536	1	41	2.0	1	30.0	0.0	0	0	0	228.0	113.0	82.5
4089	0	64	1.0	0	0.0	0.0	0	1	0	232.0	149.5	84.0
...
3444	0	36	1.0	1	5.0	0.0	0	1	0	222.0	147.0	94.0
466	0	57	3.0	1	15.0	0.0	0	0	0	250.0	125.0	74.0
3092	0	60	2.0	0	0.0	0.0	0	1	0	298.0	133.0	89.0
3772	1	39	2.0	1	10.0	0.0	0	0	0	215.0	102.0	64.5
860	0	35	2.0	0	0.0	0.0	0	0	0	248.0	107.0	73.0

3390 rows × 15 columns

```
In [45]: y_train
```

```
Out[45]:
```

3252	0
3946	0
1261	0
2536	0
4089	0
...	..
3444	0
466	0
3092	0
3772	0
860	0

Name: TenYearCHD, Length: 3390, dtype: int64

```
In [47]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression().fit(x_train,y_train)
model.score(x_train, y_train)
```

```
Out[47]: 0.8495575221238938
```

```
In [ ]:
```